

```
In [7]: from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
```

```
In [8]: batch_size = 128
num_classes = 10
epochs = 10

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if K.image_data_format() == 'channels_first':
    #https://machinelearningmastery.com/a-gentle-introduction-to-channels-first-and-channels-last-image-formats-for-deep-learning/
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>)
11490434/11490434 [=====] - 3s 0us/step

```
In [9]: x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255 #normalizing
x_test /= 255 #normalizing
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

```
In [14]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time

# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [15]: model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(Conv2D(64, (3, 3), activation='relu'))
#model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
conv2d_1 (Conv2D)	(None, 24, 24, 128)	36992
conv2d_2 (Conv2D)	(None, 22, 22, 64)	73792
flatten (Flatten)	(None, 30976)	0
dense (Dense)	(None, 10)	309770
=====		
Total params: 420,874		
Trainable params: 420,874		
Non-trainable params: 0		


```
In [16]: model.compile(loss=keras.losses.categorical_crossentropy,
                    optimizer=keras.optimizers.Adadelta(),
                    metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
Epoch 1/10
469/469 [=====] - 257s 542ms/step - loss: 2.2711 - accuracy: 0.3253 - val_loss: 2.2329 - val_accuracy:
0.4806
Epoch 2/10
469/469 [=====] - 240s 511ms/step - loss: 2.1498 - accuracy: 0.5391 - val_loss: 2.0008 - val_accuracy:
0.6362
Epoch 3/10
469/469 [=====] - 240s 513ms/step - loss: 1.6109 - accuracy: 0.7204 - val_loss: 1.0603 - val_accuracy:
0.7987
Epoch 4/10
469/469 [=====] - 242s 516ms/step - loss: 0.7814 - accuracy: 0.8094 - val_loss: 0.5816 - val_accuracy:
0.8428
Epoch 5/10
469/469 [=====] - 253s 540ms/step - loss: 0.5531 - accuracy: 0.8418 - val_loss: 0.4835 - val_accuracy:
0.8627
Epoch 6/10
469/469 [=====] - 278s 592ms/step - loss: 0.4891 - accuracy: 0.8579 - val_loss: 0.4437 - val_accuracy:
0.8719
Epoch 7/10
469/469 [=====] - 273s 583ms/step - loss: 0.4569 - accuracy: 0.8667 - val_loss: 0.4202 - val_accuracy:
0.8787
Epoch 8/10
469/469 [=====] - 245s 523ms/step - loss: 0.4359 - accuracy: 0.8733 - val_loss: 0.4036 - val_accuracy:
0.8840
Epoch 9/10
469/469 [=====] - 265s 566ms/step - loss: 0.4202 - accuracy: 0.8778 - val_loss: 0.3905 - val_accuracy:
0.8888
Epoch 10/10
469/469 [=====] - 256s 546ms/step - loss: 0.4079 - accuracy: 0.8817 - val_loss: 0.3803 - val_accuracy:
0.8919
Test loss: 0.38033363223075867
Test accuracy: 0.8919000029563904
```

```
In [17]: import matplotlib.pyplot as plt
%matplotlib inline
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# List of epoch numbers
x = list(range(1,epochs+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a List of Length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.38033363223075867
Test accuracy: 0.8919000029563904



```
In [18]: model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
activation='relu',
input_shape=input_shape))
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 26, 26, 32)	320
conv2d_4 (Conv2D)	(None, 24, 24, 128)	36992
conv2d_5 (Conv2D)	(None, 22, 22, 64)	73792
conv2d_6 (Conv2D)	(None, 20, 20, 32)	18464
flatten_1 (Flatten)	(None, 12800)	0
dense_1 (Dense)	(None, 10)	128010
=====		
Total params: 257,578		
Trainable params: 257,578		
Non-trainable params: 0		

```
In [19]: model.compile(loss=keras.losses.categorical_crossentropy,
                    optimizer=keras.optimizers.Adam(),
                    metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```


```
Epoch 1/10
469/469 [=====] - 376s 770ms/step - loss: 0.1581 - accuracy: 0.9538 - val_loss: 0.0465 - val_accuracy:
0.9837
Epoch 2/10
469/469 [=====] - 308s 656ms/step - loss: 0.0454 - accuracy: 0.9861 - val_loss: 0.0430 - val_accuracy:
0.9863
Epoch 3/10
469/469 [=====] - 301s 641ms/step - loss: 0.0309 - accuracy: 0.9902 - val_loss: 0.0410 - val_accuracy:
0.9870
Epoch 4/10
469/469 [=====] - 300s 639ms/step - loss: 0.0230 - accuracy: 0.9924 - val_loss: 0.0392 - val_accuracy:
0.9892
Epoch 5/10
469/469 [=====] - 312s 665ms/step - loss: 0.0181 - accuracy: 0.9942 - val_loss: 0.0401 - val_accuracy:
0.9887
Epoch 6/10
469/469 [=====] - 402s 856ms/step - loss: 0.0140 - accuracy: 0.9952 - val_loss: 0.0368 - val_accuracy:
0.9888
Epoch 7/10
469/469 [=====] - 351s 748ms/step - loss: 0.0107 - accuracy: 0.9963 - val_loss: 0.0372 - val_accuracy:
0.9894
Epoch 8/10
469/469 [=====] - 2942s 6s/step - loss: 0.0100 - accuracy: 0.9966 - val_loss: 0.0487 - val_accuracy:
0.9881
Epoch 9/10
469/469 [=====] - 318s 677ms/step - loss: 0.0067 - accuracy: 0.9978 - val_loss: 0.0654 - val_accuracy:
0.9855
Epoch 10/10
469/469 [=====] - 313s 667ms/step - loss: 0.0066 - accuracy: 0.9977 - val_loss: 0.0568 - val_accuracy:
0.9867
Test loss: 0.05681750550866127
Test accuracy: 0.9866999983787537
```

```

In [20]: import matplotlib.pyplot as plt
%matplotlib inline
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# List of epoch numbers
x = list(range(1,epochs+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a List of Length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.05681750550866127
 Test accuracy: 0.9866999983787537



```
In [22]: from keras.layers import BatchNormalization
from keras.layers import Dropout

model = Sequential()
model.add(Conv2D(32, kernel_size=(5,5),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Conv2D(64, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_7 (Conv2D)	(None, 24, 24, 32)	832
conv2d_8 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d (MaxPooling2D)	(None, 10, 10, 64)	0
batch_normalization (Batch Normalization)	(None, 10, 10, 64)	256
dropout (Dropout)	(None, 10, 10, 64)	0
conv2d_9 (Conv2D)	(None, 6, 6, 64)	102464
max_pooling2d_1 (MaxPooling2D)	(None, 3, 3, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 3, 3, 64)	256
dropout_1 (Dropout)	(None, 3, 3, 64)	0
flatten_2 (Flatten)	(None, 576)	0
dense_2 (Dense)	(None, 10)	5770
=====		
Total params: 160,842		
Trainable params: 160,586		
Non-trainable params: 256		

```
In [23]: model.compile(loss=keras.losses.categorical_crossentropy,
                    optimizer=keras.optimizers.Adadelta(),
                    metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Epoch 1/10
469/469 [=====] - 557s 789ms/step - loss: 3.2142 - accuracy: 0.1488 - val_loss: 2.3515 - val_accuracy: 0.2316
Epoch 2/10
469/469 [=====] - 139s 297ms/step - loss: 2.6254 - accuracy: 0.2350 - val_loss: 1.5569 - val_accuracy: 0.4987
Epoch 3/10
469/469 [=====] - 136s 291ms/step - loss: 2.2328 - accuracy: 0.3181 - val_loss: 1.2712 - val_accuracy: 0.5992
Epoch 4/10
469/469 [=====] - 130s 277ms/step - loss: 1.9420 - accuracy: 0.3928 - val_loss: 1.0817 - val_accuracy: 0.6659
Epoch 5/10
469/469 [=====] - 125s 267ms/step - loss: 1.7272 - accuracy: 0.4520 - val_loss: 0.9391 - val_accuracy: 0.7147
Epoch 6/10
469/469 [=====] - 129s 274ms/step - loss: 1.5509 - accuracy: 0.5020 - val_loss: 0.8337 - val_accuracy: 0.7529
Epoch 7/10
469/469 [=====] - 130s 277ms/step - loss: 1.4179 - accuracy: 0.5419 - val_loss: 0.7492 - val_accuracy: 0.7830
Epoch 8/10
469/469 [=====] - 144s 307ms/step - loss: 1.3182 - accuracy: 0.5712 - val_loss: 0.6826 - val_accuracy: 0.8048
Epoch 9/10
469/469 [=====] - 147s 312ms/step - loss: 1.2064 - accuracy: 0.6068 - val_loss: 0.6264 - val_accuracy: 0.8224
Epoch 10/10
469/469 [=====] - 137s 291ms/step - loss: 1.1318 - accuracy: 0.6298 - val_loss: 0.5796 - val_accuracy: 0.8364
Test loss: 0.5795929431915283
Test accuracy: 0.8363999724388123


```

In [24]: import matplotlib.pyplot as plt
%matplotlib inline
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# List of epoch numbers
x = list(range(1,epochs+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a List of Length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.5795929431915283

Test accuracy: 0.8363999724388123




```
In [25]: model = Sequential()
model.add(Conv2D(32, kernel_size=(2,2),
                strides=(2,2),
                padding='same',
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (2, 2), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Conv2D(64, (2, 2), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_10 (Conv2D)	(None, 14, 14, 32)	160
conv2d_11 (Conv2D)	(None, 13, 13, 64)	8256
max_pooling2d_2 (MaxPooling 2D)	(None, 6, 6, 64)	0
batch_normalization_2 (Batch Normalization)	(None, 6, 6, 64)	256
dropout_2 (Dropout)	(None, 6, 6, 64)	0
conv2d_12 (Conv2D)	(None, 5, 5, 64)	16448
max_pooling2d_3 (MaxPooling 2D)	(None, 2, 2, 64)	0
batch_normalization_3 (Batch Normalization)	(None, 2, 2, 64)	256
dropout_3 (Dropout)	(None, 2, 2, 64)	0
flatten_3 (Flatten)	(None, 256)	0
dense_3 (Dense)	(None, 10)	2570
=====		
Total params: 27,946		
Trainable params: 27,690		
Non-trainable params: 256		

```
In [26]: model.compile(loss=keras.losses.categorical_crossentropy,
                    optimizer=keras.optimizers.Adadelta(),
                    metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Epoch 1/10
469/469 [=====] - 85s 132ms/step - loss: 3.6439 - accuracy: 0.1099 - val_loss: 3.5213 - val_accuracy: 0.0902
Epoch 2/10
469/469 [=====] - 31s 66ms/step - loss: 3.5594 - accuracy: 0.1156 - val_loss: 2.7564 - val_accuracy: 0.0813
Epoch 3/10
469/469 [=====] - 30s 64ms/step - loss: 3.4622 - accuracy: 0.1222 - val_loss: 2.6716 - val_accuracy: 0.1039
Epoch 4/10
469/469 [=====] - 31s 66ms/step - loss: 3.3829 - accuracy: 0.1291 - val_loss: 2.6045 - val_accuracy: 0.1298
Epoch 5/10
469/469 [=====] - 30s 65ms/step - loss: 3.3127 - accuracy: 0.1352 - val_loss: 2.5411 - val_accuracy: 0.1499
Epoch 6/10
469/469 [=====] - 31s 66ms/step - loss: 3.2227 - accuracy: 0.1437 - val_loss: 2.4795 - val_accuracy: 0.1694
Epoch 7/10
469/469 [=====] - 29s 63ms/step - loss: 3.1713 - accuracy: 0.1508 - val_loss: 2.4207 - val_accuracy: 0.1817
Epoch 8/10
469/469 [=====] - 30s 65ms/step - loss: 3.0852 - accuracy: 0.1587 - val_loss: 2.3621 - val_accuracy: 0.1933
Epoch 9/10
469/469 [=====] - 25s 54ms/step - loss: 3.0322 - accuracy: 0.1666 - val_loss: 2.3047 - val_accuracy: 0.2062
Epoch 10/10
469/469 [=====] - 26s 55ms/step - loss: 2.9765 - accuracy: 0.1759 - val_loss: 2.2495 - val_accuracy: 0.2164
Test loss: 2.249525308609009
Test accuracy: 0.2163999747276306


```

In [28]: import matplotlib.pyplot as plt
%matplotlib inline
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# List of epoch numbers
x = list(range(1,epochs+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a List of Length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 2.249525308609009

Test accuracy: 0.21639999747276306



```
In [29]: model = Sequential()
model.add(Conv2D(32, kernel_size=(7,7),
                padding='valid',
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(128, (7, 7), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Conv2D(64, (7, 7), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Flatten())
model.add(Dense(num_classes, activation='softmax'))

model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
=====		
conv2d_13 (Conv2D)	(None, 22, 22, 32)	1600
conv2d_14 (Conv2D)	(None, 16, 16, 128)	200832
max_pooling2d_4 (MaxPooling 2D)	(None, 8, 8, 128)	0
batch_normalization_4 (Batch Normalization)	(None, 8, 8, 128)	512
dropout_4 (Dropout)	(None, 8, 8, 128)	0
conv2d_15 (Conv2D)	(None, 2, 2, 64)	401472
max_pooling2d_5 (MaxPooling 2D)	(None, 1, 1, 64)	0
batch_normalization_5 (Batch Normalization)	(None, 1, 1, 64)	256
dropout_5 (Dropout)	(None, 1, 1, 64)	0
flatten_4 (Flatten)	(None, 64)	0
dense_4 (Dense)	(None, 10)	650
=====		
Total params: 605,322		
Trainable params: 604,938		
Non-trainable params: 384		

```
In [30]: model.compile(loss=keras.losses.categorical_crossentropy,
                    optimizer=keras.optimizers.Adam(),
                    metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Epoch 1/10
469/469 [=====] - 276s 574ms/step - loss: 0.1735 - accuracy: 0.9499 - val_loss: 0.1095 - val_accuracy: 0.9825
Epoch 2/10
469/469 [=====] - 267s 569ms/step - loss: 0.0640 - accuracy: 0.9810 - val_loss: 0.0273 - val_accuracy: 0.9910
Epoch 3/10
469/469 [=====] - 252s 537ms/step - loss: 0.0507 - accuracy: 0.9852 - val_loss: 0.0245 - val_accuracy: 0.9914
Epoch 4/10
469/469 [=====] - 256s 546ms/step - loss: 0.0424 - accuracy: 0.9873 - val_loss: 0.0242 - val_accuracy: 0.9922
Epoch 5/10
469/469 [=====] - 258s 551ms/step - loss: 0.0361 - accuracy: 0.9892 - val_loss: 0.0236 - val_accuracy: 0.9921
Epoch 6/10
469/469 [=====] - 266s 567ms/step - loss: 0.0317 - accuracy: 0.9902 - val_loss: 0.0244 - val_accuracy: 0.9923
Epoch 7/10
469/469 [=====] - 282s 602ms/step - loss: 0.0289 - accuracy: 0.9911 - val_loss: 0.0377 - val_accuracy: 0.9886
Epoch 8/10
469/469 [=====] - 307s 654ms/step - loss: 0.0262 - accuracy: 0.9918 - val_loss: 0.0191 - val_accuracy: 0.9934
Epoch 9/10
469/469 [=====] - 293s 625ms/step - loss: 0.0266 - accuracy: 0.9917 - val_loss: 0.0204 - val_accuracy: 0.9936
Epoch 10/10
469/469 [=====] - 303s 646ms/step - loss: 0.0234 - accuracy: 0.9929 - val_loss: 0.0201 - val_accuracy: 0.9935
Test loss: 0.020070811733603477
Test accuracy: 0.9934999942779541


```

In [31]: import matplotlib.pyplot as plt
%matplotlib inline
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
# List of epoch numbers
x = list(range(1,epochs+1))
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a List of Length equal to number of epochs
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.020070811733603477

Test accuracy: 0.9934999942779541



```

In [40]: from prettytable import PrettyTable
x=PrettyTable()
x.field_names=["Model", "#Hidden Layers", "Kernel-Size", "MaxPooling", "Dropout/BatchNormalization", "Optimizer", "Activation", "Accuracy"]
x.add_row(["1.", "2(128-64)", "3X3", "False", "False", "Adadelata", "ReLu", "0.989"])
x.add_row(["2.", "3(128-64-32)", "3X3", "False", "False", "Adam", "ReLu", "0.988"])
x.add_row(["3.", "2(64-64)", "5X5", "2X2", "True", "Adadelata", "ReLu", "0.993"])
x.add_row(["4.", "2(64-64)", "2X2(s=2,p='same')", "2X2", "True", "Adadelata", "ReLu", "0.983"])
x.add_row(["5.", "2(128-64)", "7X7(s=2,p='valid')", "2X2", "True", "Adam", "ReLu", "0.992"])
print(x)

```

Model	#Hidden Layers	Kernel-Size	MaxPooling	Dropout/BatchNormalization	Optimizer	Activation	Accuracy
1.	2(128-64)	3X3	False	False	Adadelata	ReLu	0.989
2.	3(128-64-32)	3X3	False	False	Adam	ReLu	0.988
3.	2(64-64)	5X5	2X2	True	Adadelata	ReLu	0.993
4.	2(64-64)	2X2(s=2,p='same')	2X2	True	Adadelata	ReLu	0.983
5.	2(128-64)	7X7(s=2,p='valid')	2X2	True	Adam	ReLu	0.992